

Internal Policy Guidelines

Data Governance, Content Standards & AI Tool Usage — Version 3.2

Document Type:	Policy & Compliance Document
Prepared:	March 1, 2025
Authors:	Legal, Compliance & Data Governance Office
Version:	3.2
Classification:	CONFIDENTIAL — Internal Use Only

Table of Contents

1. Scope & Purpose
2. Data Classification & Access Control
3. Viewer Data Privacy Policy
4. Content Quality Standards
5. AI & Analytics Tool Usage Policy
6. Marketing & Advertising Standards
7. Incident Reporting
8. Policy Enforcement & Disciplinary Action

1. Scope & Purpose

These guidelines govern all employees, contractors, and third-party partners of StreamVault Entertainment who access, process, or act upon internal business data, viewer information, content assets, or AI-powered analytical tools. Compliance is mandatory. Violations are subject to disciplinary action up to and including termination and legal proceedings.

This document supersedes Policy Guidelines v3.1 (October 2024) and should be read in conjunction with the StreamVault Employee Code of Conduct (HR-POL-001) and the Information Security Policy (IS-POL-007).

2. Data Classification & Access Control

All internal data is classified into four tiers. Access is granted on a need-to-know basis and must be approved by the data owner of each tier.

Tier	Classification	Examples	Access Requires
T1	Public	Press releases, published film metadata	No approval — publicly available
T2	Internal	Aggregate genre stats, general KPIs	Employee login
T3	Confidential	Viewer PII, financial data, salary info	Manager approval + audit log
T4	Restricted	Raw viewer data, AI model weights, legal	VP sign-off + Legal review

The AI Insights Assistant (Secure AI system) is authorised to access T2 and T3 data only via approved backend tool calls. Direct T4 raw data access by the AI layer is prohibited. All AI queries are logged in the audit trail.

3. Viewer Data Privacy Policy

StreamVault processes viewer data in compliance with the **General Data Protection Regulation (GDPR)** for EU/UK users, the **Digital Personal Data Protection Act 2023 (DPDPA)** for Indian users, and the **California Consumer Privacy Act (CCPA)** for California-based users. The following rules apply universally:

- Viewer PII (name, email, age, location) must never be included verbatim in AI prompt context or response outputs. Use anonymised viewer_id references only.
- Watch history and review data may be used in aggregate for analytics but may not be associated with an identifiable individual without explicit viewer consent.
- Viewer data retention periods: Active account data — retained for the duration of subscription plus 2 years. Inactive accounts — purged after 3 years of inactivity.
- Cross-border data transfers: Viewer data for Indian users must reside on India-region servers (AWS ap-south-1 or equivalent). Transfers to EU/US servers require DPA sign-off.
- Right to erasure: Viewer deletion requests must be fulfilled within 30 days. The data engineering team maintains a deletion workflow documented in runbook DE-RB-009.

4. Content Quality Standards

All commissioned content must meet minimum quality gates before receiving production greenlight. These thresholds are based on historical correlation between pre-production quality indicators and post-release platform metrics.

Gate	Criteria	Minimum Threshold	Consequence if Failed
Script Review	Internal story dept score	>= 7.5 / 10	Script revision — 4-week hold
Pilot Audience	Focus group completion rate	>= 68%	Concept rework required
Budget Alignment	Production budget vs P&L; model	<= 22% of rev	Finance approval required
Genre Fit	Hist. platform performance of genre	!= PAUSE	CCO sign-off mandatory
DEI Assessment	Inclusion checklist score	>= 80 / 100	Revision required

Comedy Audit Mandate (effective Q1 2025): All four Comedy titles released in 2024-Q4 through 2025-Q1 failed the Script Review gate on retrospective evaluation (average score: 5.8/10). Comedy commissioning is suspended until a root-cause analysis report is submitted to the CCO no later than June 30, 2025.

5. AI & Analytics Tool Usage Policy

The **StreamVault AI Insights Assistant** is an internal tool authorised for use by Analytics, Strategy, Content, and Marketing teams. The following rules govern its use:

- **Permitted uses:** Querying aggregate business metrics, trend analysis, content performance comparison, marketing ROI questions, and leadership recommendation generation.
- **Prohibited uses:** Querying individual viewer PII, generating external-facing content without editorial review, automating financial transactions, or accessing T4 restricted data.
- **Output validation:** AI-generated reports and recommendations must be reviewed by a qualified analyst before distribution to leadership. AI outputs are advisory only.
- **Audit logging:** Every query to the AI Insights Assistant is logged with: user ID, timestamp, query text, tools invoked, and data sources accessed.
- **Prompt injection protection:** The AI system must validate all user inputs against an allowlist of query patterns. Attempts to override system instructions must be blocked and flagged for Security review.
- **Model version control:** Only approved model versions may be deployed. Model changes require a Security Review Board approval and a 72-hour staged rollout.

6. Marketing & Advertising Standards

All marketing campaigns for StreamVault content must comply with applicable advertising standards in the markets where they are served, including but not limited to the Advertising Standards Council of India (ASCI) Guidelines, the UK Advertising Standards Authority (ASA) Code, and the US Federal Trade Commission (FTC) guidelines on endorsements and testimonials.

Influencer and creator partnerships must include clear disclosure of the paid commercial relationship. Undisclosed sponsored content is a legal and reputational risk and will result in immediate campaign suspension.

7. Incident Reporting

Any suspected breach of these policies — including unauthorised data access, AI system misuse, content standard violations, or marketing compliance failures — must be reported within **24 hours** of discovery to the Compliance Hotline (compliance@streamvault.internal) or via the anonymous Ethics Reporting Portal.

Incident Type	Report To	Response SLA
Data breach / PII exposure	CISO + Legal + DPO	2 hours
AI system misuse	Head of AI Governance	4 hours
Content quality failure	Chief Content Officer	24 hours
Marketing compliance breach	VP Marketing + Legal	8 hours
Financial data unauthorised access	CFO + Internal Audit	4 hours

8. Policy Enforcement & Disciplinary Action

Violations of these guidelines are categorised by severity. Repeat violations escalate automatically to the next severity tier regardless of initial category.

Severity	Examples	Consequence
Minor	Accessing internal data outside normal workflow	Verbal warning + mandatory training
Moderate	Sharing T3 data without approval	Written warning + access suspension
Serious	Exposing viewer PII; bypassing AI safety controls	Suspension pending investigation
Critical	Intentional data breach; regulatory non-compliance	Termination + potential legal action

Approved by: Chief Compliance Officer | Effective Date: March 1, 2025 | Next Review: September 1, 2025 | Document Owner: Legal & Compliance Office | Policy ID: COMP-POL-012-v3.2